

✓ Carga y análisis del dataset

Instalar librerías necesarias (si no están instaladas)

```
!pip install pandas openpyxl matplotlib seaborn --quiet
!pip install datasets --quiet
!pip install numpy==1.26.4 --quiet # (versión anterior)
```

Después de ejecutar estos comandos es necesario reiniciar el kernel



```
61.0/61.0 kB 4.0 MB/s eta 0:00:00
18.3/18.3 MB 46.6 MB/s eta 0:00:00
ERROR: pip's dependency resolver does not currently take into account all the packages that are ins
thinc 8.3.6 requires numpy<3.0.0,>=2.0.0, but you have numpy 1.26.4 which is incompatible.
```

Importar librerías

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
from transformers import BertTokenizer, BertForSequenceClassification, Trainer, TrainingArguments
from transformers import EarlyStoppingCallback
from datasets import Dataset
import torch
```

Carga del dataset y normalización

```
ruta = "/content/Dataset_Completo.xlsx"
df = pd.read_excel(ruta)

# Limpieza y normalización de la columna 'Category'
df['Category'] = df['Category'].astype(str).str.strip().str.lower()

# Mapeo de etiquetas: 'true' → 0 y 'fake' → 1
df['label'] = df['Category'].map({'true': 0, 'fake': 1})

# Verificación
print("Valores únicos en 'Category':", df['Category'].unique())
print("Valores únicos en 'label':", df['label'].unique())
print("\nDistribución de clases:")
print(df['label'].value_counts())
```



```
Valores únicos en 'Category': ['fake' 'true']
Valores únicos en 'label': [1 0]
```

```
Distribución de clases:
label
0      777
1      766
Name: count, dtype: int64
```

✓ Entrenamiento de BERT con el texto completo

```
from sklearn.model_selection import train_test_split

# 1. Usamos 'text', 'topic', 'source' y 'headline' como entrada. El campo LINK no lo utilizamos para el
df_texto = df[['Text', 'Topic', 'Source', 'Headline', 'label']].dropna().copy()

# Renombramos todas las columnas
df_texto.rename(columns={
    'Text': 'text',
    'Topic': 'topic',
    'Source': 'source',
    'Headline': 'headline'
}, inplace=True)

# 2. División del dataset en entrenamiento y prueba
train_df, test_df = train_test_split(df_texto, test_size=0.2, stratify=df_texto['label'], random_state=

# 3. Conversión a datasets de Hugging Face
train_dataset = Dataset.from_pandas(train_df)
test_dataset = Dataset.from_pandas(test_df)

# 4. Tokenizador y tokenización
tokenizer = BertTokenizer.from_pretrained("bert-base-multilingual-cased")

def tokenize_function(examples):
    return tokenizer(examples["text"], padding="max_length", truncation=True, max_length=512)

tokenized_train = train_dataset.map(tokenize_function, batched=True)
tokenized_test = test_dataset.map(tokenize_function, batched=True)

# 5. Formateo
tokenized_train.set_format("torch", columns=["input_ids", "attention_mask", "label"])
tokenized_test.set_format("torch", columns=["input_ids", "attention_mask", "label"])

# 6. Modelo multilingüe
model = BertForSequenceClassification.from_pretrained("bert-base-multilingual-cased", num_labels=2)

# 7. Configuración de entrenamiento
training_args = TrainingArguments(
    output_dir="./results_multilingual",
    num_train_epochs=10,
    per_device_train_batch_size=16,
    per_device_eval_batch_size=64,
    eval_strategy="epoch", # Para evaluar cada época
    save_strategy="epoch",
    load_best_model_at_end=True,
    metric_for_best_model="eval_loss",
    greater_is_better=False,
    logging_steps=10,
    report_to="none"
)

# 8 Preparación metricas entrenamiento
from sklearn.metrics import accuracy_score

def compute_metrics(eval_pred):
    logits, labels = eval_pred
    preds = logits.argmax(axis=-1)
    acc = accuracy_score(labels, preds)
```

```

    return {"accuracy": acc}

# 9. Entrenador
trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=tokenized_train,
    eval_dataset=tokenized_test,
    compute_metrics=compute_metrics,
    callbacks=[EarlyStoppingCallback(early_stopping_patience=2)], # Espera 2 épocas sin mejorar
)

# 9. Entrenar el modelo
trainer.train()

```

➡ /usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94: UserWarning:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings tab (<https://huggingface.co/settings/tokens>)
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access public models or datasets

warnings.warn(
tokenizer_config.json: 100% 49.0/49.0 [00:00<00:00, 1.85kB/s]
vocab.txt: 100% 996k/996k [00:00<00:00, 6.62MB/s]
tokenizer.json: 100% 1.96M/1.96M [00:00<00:00, 7.33MB/s]
config.json: 100% 625/625 [00:00<00:00, 24.8kB/s]
Map: 100% 1234/1234 [00:19<00:00, 66.80 examples/s]
Map: 100% 309/309 [00:02<00:00, 153.38 examples/s]
Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed. Falling back to regular HTTP.
WARNING:huggingface_hub.file_download:Xet Storage is enabled for this repo, but the 'hf_xet' package is not installed.
model.safetensors: 100% 714M/714M [00:03<00:00, 260MB/s]
Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-multilingual-uncased. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and generation.

[234/780 09:22 < 22:03, 0.41 it/s, Epoch 3/10]

Epoch	Training Loss	Validation Loss	Accuracy
1	0.556600	0.449987	0.796117
2	0.444900	0.512026	0.773463
3	0.273500	0.471875	0.809061

TrainOutput(global_step=234, training_loss=0.5157423651116526, metrics={'train_runtime': 564.824, 'train_samples_per_second': 21.848, 'train_steps_per_second': 1.381, 'total_flos': 974037126942720.0, 'train_loss': 0.5157423651116526, 'epoch': 3.0})

Guardar el modelo entrenado

```

# Guardar modelo y tokenizer tras el entrenamiento
trainer.save_model("./modelo_bert_multilingual")
tokenizer.save_pretrained("./modelo_bert_multilingual")

```

➡ ('./modelo_bert_multilingual/tokenizer_config.json',
'./modelo_bert_multilingual/special_tokens_map.json',
'./modelo_bert_multilingual/vocab.txt',
'./modelo_bert_multilingual/added_tokens.json')

✓ Evaluación del modelo con texto completo

```
from sklearn.metrics import (
    accuracy_score, precision_score, recall_score, f1_score,
    classification_report, confusion_matrix, roc_auc_score, roc_curve
)

# 1. Obtener predicciones
predictions = trainer.predict(tokenized_test)
preds = np.argmax(predictions.predictions, axis=1)
labels = predictions.label_ids

# 2. Métricas principales
accuracy = accuracy_score(labels, preds)
precision = precision_score(labels, preds)
recall = recall_score(labels, preds)
f1 = f1_score(labels, preds)

print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-score: {f1:.4f}")

# 3. Reporte completo por clase
print("\nReporte completo:")
print(classification_report(labels, preds, target_names=["True", "Fake"]))

# 4. Matriz de confusión
cm = confusion_matrix(labels, preds)
plt.figure(figsize=(6, 4))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=["True", "Fake"], yticklabels=["True", "Fake"])
plt.xlabel('Predicción')
plt.ylabel('Real')
plt.title('Matriz de Confusión - Texto completo (BERT)')
plt.show()

# 5. Curva ROC y AUC
probs = torch.nn.functional.softmax(torch.tensor(predictions.predictions), dim=1)[ :, 1].numpy()
auc = roc_auc_score(labels, probs)
fpr, tpr, _ = roc_curve(labels, probs)

plt.figure(figsize=(7, 5))
plt.plot(fpr, tpr, label=f'ROC curve (AUC = {auc:.4f})')
plt.plot([0, 1], [0, 1], 'k--', label='Clasificador aleatorio')
plt.xlabel('FPR')
plt.ylabel('TPR')
plt.title('Curva ROC - Modelo con texto completo (BERT)')
plt.legend(loc='lower right')
plt.grid(True, linestyle='--', alpha=0.6)
plt.tight_layout()
plt.show()
```



Accuracy: 0.7961
Precision: 0.8409
Recall: 0.7255
F1-score: 0.7789

Reporte completo:

	precision	recall	f1-score	support
True	0.76	0.87	0.81	156
Fake	0.84	0.73	0.78	153
accuracy			0.80	309
macro avg	0.80	0.80	0.79	309
weighted avg	0.80	0.80	0.80	309

